



| |
|---|
| Experiment No. 3 |
| Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model |
| Date of Performance:07-08-23 |
| Date of Submission:17-08-23 |

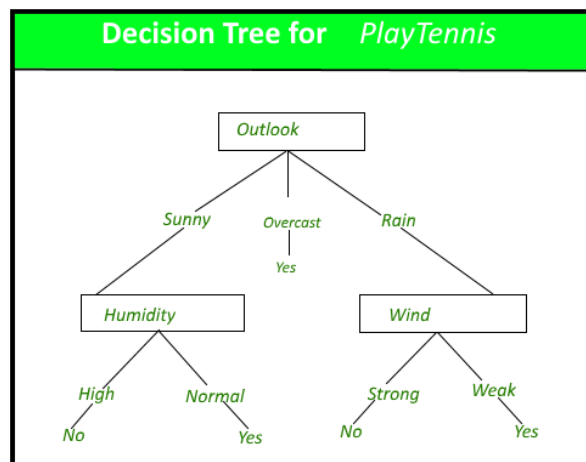


Aim: Apply Decision Tree Algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: To perform various feature engineering tasks, apply Decision Tree Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score. Improve the performance by performing different data engineering and feature engineering tasks.

Theory:

Decision Tree is the most powerful and popular tool for classification and prediction. A Decision tree is a flowchart-like tree structure, where each internal node denotes a test on an attribute, each branch represents an outcome of the test, and each leaf node (terminal node) holds a class label.



Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras,



Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

Code:

Conclusion:

The Decision Tree model demonstrated good performance on the Adult Census Income Dataset. It managed categorical attributes using one-hot encoding and performed necessary data preprocessing, such as handling missing values, dropping irrelevant tables, and separating columns, to enhance the model's effectiveness.

Hyperparameter tuning is a critical process for enhancing the Decision Tree model's performance, as it allows the control over the model's complexity by setting some limits on the parameters. To take performance to the next level, we need to improve the model by adjusting hyperparameters like max depth, min samples split, etc., using methods like Grid Search or Random Search.

Accuracy: Achieved an accuracy of 0.85, indicating that around 85% of predictions were correct.

Confusion Matrix: True positive = 9860, False positive = 481, False negative = 1823, True negative = 1646

Precision: Precision of 0.84 suggests that among the instances predicted as 0, about 0.77 are predicted for 1.

Recall: Recall of 0.95 indicates that the model captured the instances for 0 and 0.47 model captured the instances for 1.

F1 Score: The F1 score of 0.90 is the mean of precision and recall for 0 and 0.59 is the mean of precision and recall for 1 in the model's performance.

```
# Import libraries
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
# To ignore warning messages
import warnings
warnings.filterwarnings('ignore')
# Adult dataset path
adult_dataset_path = "/content/adult.csv"

# Function for loading adult dataset
def load_adult_data(adult_path=adult_dataset_path):
    csv_path = os.path.join(adult_path)
    return pd.read_csv(csv_path)

# Calling load adult function and assigning to a new variable df
df = load_adult_data()
# load top 3 rows values from adult dataset
df.head(3)
```

| | age | workclass | fnlwgt | education | educational-num | marital-status | occupation | relationship |
|---|-----|-----------|--------|------------|-----------------|--------------------|-------------------|--------------|
| 0 | 25 | Private | 226802 | 11th | 7 | Never-married | Machine-op-inspct | Own-child |
| 1 | 38 | Private | 89814 | HS-grad | 9 | Married-civ-spouse | Farming-fishing | Husband |
| 2 | 28 | Local-gov | 336951 | Assoc-acdm | 12 | Married-civ-spouse | Protective-serv | Husband |

```
print ("Rows : " ,df.shape[0])
print ("Columns : " ,df.shape[1])
print ("\nFeatures : \n" ,df.columns.tolist())
print ("\nMissing values : " , df.isnull().sum().values.sum())
print ("\nUnique values : \n",df.nunique())
```

```
Rows : 48842
Columns : 15
```

```
Features :
['age', 'workclass', 'fnlwgt', 'education', 'educational-num', 'marital-status', 'occupation', 'relationship', 'r
```

```
Missing values : 0
```

```
Unique values :
age                74
workclass           9
fnlwgt            28523
education          16
educational-num    16
marital-status      7
occupation         15
relationship        6
race                5
gender              2
capital-gain       123
```

```
capital-loss      99
hours-per-week    96
native-country    42
income            2
dtype: int64
```

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48842 entries, 0 to 48841
Data columns (total 15 columns):
#   Column              Non-Null Count  Dtype
---  -
0   age                 48842 non-null  int64
1   workclass           48842 non-null  object
2   fnlwgt              48842 non-null  int64
3   education           48842 non-null  object
4   educational-num     48842 non-null  int64
5   marital-status      48842 non-null  object
6   occupation          48842 non-null  object
7   relationship        48842 non-null  object
8   race                48842 non-null  object
9   gender              48842 non-null  object
10  capital-gain         48842 non-null  int64
11  capital-loss         48842 non-null  int64
12  hours-per-week       48842 non-null  int64
13  native-country      48842 non-null  object
14  income              48842 non-null  object
dtypes: int64(6), object(9)
memory usage: 5.6+ MB
```

df.describe()

| | age | fnlwgt | educational-num | capital-gain | capital-loss | hours-per-week |
|-------|--------------|--------------|-----------------|--------------|--------------|----------------|
| count | 48842.000000 | 4.884200e+04 | 48842.000000 | 48842.000000 | 48842.000000 | 48842.000000 |
| mean | 38.643585 | 1.896641e+05 | 10.078089 | 1079.067626 | 87.502314 | 40.422382 |
| std | 13.710510 | 1.056040e+05 | 2.570973 | 7452.019058 | 403.004552 | 12.391444 |
| min | 17.000000 | 1.228500e+04 | 1.000000 | 0.000000 | 0.000000 | 1.000000 |
| 25% | 28.000000 | 1.175505e+05 | 9.000000 | 0.000000 | 0.000000 | 40.000000 |
| 50% | 37.000000 | 1.781445e+05 | 10.000000 | 0.000000 | 0.000000 | 40.000000 |
| 75% | 48.000000 | 2.376420e+05 | 12.000000 | 0.000000 | 0.000000 | 45.000000 |

df.head()

| | age | workclass | fnlwgt | education | educational-num | marital-status | occupation | relationship |
|---|-----|-----------|--------|-----------|-----------------|----------------|-------------------|--------------|
| 0 | 25 | Private | 226802 | 11th | 7 | Never-married | Machine-op-inspct | Own-child |

```
# checking "?" total values present in particular 'workclass' feature
```

```
df_check_missing_workclass = (df['workclass']=='?').sum()
```

```
df_check_missing_workclass
```

```
2799
```

```
spouse
```

```
# checking "?" total values present in particular 'occupation' feature
```

```
df_check_missing_occupation = (df['occupation']=='?').sum()
```

```
df_check_missing_occupation
```

```
2809
```

```
# checking "?" values, how many are there in the whole dataset
```

```
df_missing = (df=='?').sum()
```

```
df_missing
```

```
age                0
workclass          2799
fnlwgt             0
education          0
educational-num    0
marital-status     0
occupation         2809
relationship       0
race               0
gender             0
capital-gain       0
capital-loss       0
hours-per-week     0
native-country     857
income             0
dtype: int64
```

```
percent_missing = (df=='?').sum() * 100/len(df)
```

```
percent_missing
```

```
age                0.000000
workclass          5.730724
fnlwgt             0.000000
education          0.000000
educational-num    0.000000
marital-status     0.000000
occupation         5.751198
relationship       0.000000
race               0.000000
gender             0.000000
capital-gain       0.000000
capital-loss       0.000000
hours-per-week     0.000000
native-country     1.754637
income             0.000000
dtype: float64
```

```
# find total number of rows which doesn't contain any missing value as '?'
```

```
df.apply(lambda x: x != '?',axis=1).sum()
```

```
age                48842
workclass          46043
fnlwgt             48842
education          48842
```

```

educational-num    48842
marital-status     48842
occupation         46033
relationship       48842
race              48842
gender            48842
capital-gain      48842
capital-loss      48842
hours-per-week    48842
native-country    47985
income            48842
dtype: int64

```

```

# dropping the rows having missing values in workclass
df = df[df['workclass'] != '?']
df.head()

```

| | age | workclass | fnlwgt | education | educational-num | marital-status | occupation | relationship |
|---|-----|-----------|--------|--------------|-----------------|--------------------|-------------------|---------------|
| 0 | 25 | Private | 226802 | 11th | 7 | Never-married | Machine-op-inspct | Own-child |
| 1 | 38 | Private | 89814 | HS-grad | 9 | Married-civ-spouse | Farming-fishing | Husband |
| 2 | 28 | Local-gov | 336951 | Assoc-acdm | 12 | Married-civ-spouse | Protective-serv | Husband |
| 3 | 44 | Private | 160323 | Some-college | 10 | Married-civ-spouse | Machine-op-inspct | Husband |
| 5 | 34 | Private | 198693 | 10th | 6 | Never-married | Other-service | Not-in-family |

```

# select all categorical variables
df_categorical = df.select_dtypes(include=['object'])
# checking whether any other column contains '?' value
df_categorical.apply(lambda x: x=='?',axis=1).sum()

```

```

workclass          0
education          0
marital-status     0
occupation        10
relationship       0
race              0
gender            0
native-country    811
income            0
dtype: int64

```

```

from sklearn import preprocessing
# encode categorical variables using label Encoder
# select all categorical variables
df_categorical = df.select_dtypes(include=['object'])
df_categorical.head()

```


| | workclass | education | marital-status | occupation | relationship | race | gender | native-country | income |
|---|-----------|-----------|----------------|-------------------|--------------|-------|--------|----------------|--------|
| 0 | Private | 11th | Never-married | Machine-op-inspct | Own-child | Black | Male | United-States | <= |

```
# apply label encoder to df_categorical
le = preprocessing.LabelEncoder()
df_categorical = df_categorical.apply(le.fit_transform)
df_categorical.head()
```

| | workclass | education | marital-status | occupation | relationship | race | gender | native-country | income |
|---|-----------|-----------|----------------|------------|--------------|------|--------|----------------|--------|
| 0 | 2 | 1 | 4 | 6 | 3 | 2 | 1 | 39 | |
| 1 | 2 | 11 | 2 | 4 | 0 | 4 | 1 | 39 | |
| 2 | 1 | 7 | 2 | 10 | 0 | 4 | 1 | 39 | |
| 3 | 2 | 15 | 2 | 6 | 0 | 2 | 1 | 39 | |

```
# Next, Concatenate df_categorical dataframe with original df (dataframe)
# first, Drop earlier duplicate columns which had categorical values
df = df.drop(df_categorical.columns,axis=1)
df = pd.concat([df,df_categorical],axis=1)
df.head()
```

| | age | fnlwgt | educational-num | capital-gain | capital-loss | hours-per-week | workclass | education | marital-status | income |
|---|-----|--------|-----------------|--------------|--------------|----------------|-----------|-----------|----------------|--------|
| 0 | 25 | 226802 | 7 | 0 | 0 | 40 | 2 | 1 | 39 | |
| 1 | 38 | 89814 | 9 | 0 | 0 | 50 | 2 | 11 | 39 | |
| 2 | 28 | 336951 | 12 | 0 | 0 | 40 | 1 | 7 | 39 | |
| 3 | 44 | 160323 | 10 | 7688 | 0 | 40 | 2 | 15 | 39 | |
| 5 | 34 | 198693 | 6 | 0 | 0 | 30 | 2 | 0 | 39 | |

```
# look at column type
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 46033 entries, 0 to 48841
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   46033 non-null  int64
1   fnlwgt                46033 non-null  int64
2   educational-num        46033 non-null  int64
3   capital-gain           46033 non-null  int64
4   capital-loss           46033 non-null  int64
5   hours-per-week         46033 non-null  int64
6   workclass              46033 non-null  int64
7   education              46033 non-null  int64
8   marital-status         46033 non-null  int64
9   occupation             46033 non-null  int64
10  relationship           46033 non-null  int64
11  race                   46033 non-null  int64
12  gender                 46033 non-null  int64
13  native-country         46033 non-null  int64
14  income                 46033 non-null  int64
dtypes: int64(15)
memory usage: 5.6 MB
```

```
# convert target variable income to categorical
df['income'] = df['income'].astype('category')
# check df info again whether everything is in right format or not
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 46033 entries, 0 to 48841
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   46033 non-null  int64
1   fnlwgt                46033 non-null  int64
2   educational-num       46033 non-null  int64
3   capital-gain          46033 non-null  int64
4   capital-loss          46033 non-null  int64
5   hours-per-week        46033 non-null  int64
6   workclass             46033 non-null  int64
7   education             46033 non-null  int64
8   marital-status        46033 non-null  int64
9   occupation            46033 non-null  int64
10  relationship          46033 non-null  int64
11  race                  46033 non-null  int64
12  gender                46033 non-null  int64
13  native-country        46033 non-null  int64
14  income                46033 non-null  category
dtypes: category(1), int64(14)
memory usage: 5.3 MB
```

```
# Importing train_test_split
from sklearn.model_selection import train_test_split
# Putting independent variables/features to X
X = df.drop('income',axis=1)
# Putting response/dependent variable/feature to y
y = df['income']
```

```
X.head(3)
```

| | age | fnlwgt | educational-num | capital-gain | capital-loss | hours-per-week | workclass | education | marital-status | occupation | relationship |
|---|-----|--------|-----------------|--------------|--------------|----------------|-----------|-----------|----------------|------------|--------------|
| 0 | 25 | 226802 | 7 | 0 | 0 | 40 | 2 | 1 | 4 | 6 | 3 |
| 1 | 38 | 89814 | 9 | 0 | 0 | 50 | 2 | 11 | 2 | 4 | 0 |
| 2 | 28 | 336951 | 12 | 0 | 0 | 40 | 1 | 7 | 2 | 10 | 0 |

```
y.head(3)
```

```
0    0
1    0
2    1
Name: income, dtype: category
Categories (2, int64): [0, 1]
```

```
# Splitting the data into train and test
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.30,random_state=99)
X_train.head()
```

| | age | fnlwgt | educational- num | capital- gain | capital- loss | hours- per- week | workclass | education | marital |
|--------------|-----|--------|---------------------|------------------|------------------|------------------------|-----------|-----------|---------|
| 29293 | 39 | 203070 | 11 | 0 | 0 | 49 | 2 | 8 | |
| 3452 | 50 | 243115 | 9 | 0 | 0 | 40 | 2 | 11 | |
| 9788 | 31 | 154227 | 10 | 0 | 0 | 40 | 2 | 15 | |

```
# Importing decision tree classifier from sklearn library
from sklearn.tree import DecisionTreeClassifier
# Fitting the decision tree with default hyperparameters, apart from
# max_depth which is 5 so that we can plot and read the tree.
dt_default = DecisionTreeClassifier(max_depth=5)
dt_default.fit(X_train,y_train)
```

```
DecisionTreeClassifier
DecisionTreeClassifier(max_depth=5)
```

```
# check the evaluation metrics of our default model
# Importing classification report and confusion matrix from sklearn metrics
from sklearn.metrics import classification_report,confusion_matrix,accuracy_score
# making predictions
y_pred_default = dt_default.predict(X_test)
# Printing classifier report after prediction
print(classification_report(y_test,y_pred_default))
```

```
precision    recall  f1-score   support

0           0.86       0.95       0.90       10341
1           0.79       0.53       0.64        3469

accuracy          0.85       13810
macro avg         0.83       0.74       0.77       13810
weighted avg      0.84       0.85       0.84       13810
```

```
# Printing confusion matrix and accuracy
print(confusion_matrix(y_test,y_pred_default))
print(accuracy_score(y_test,y_pred_default))
```

```
[[9861  480]
 [1616 1853]]
0.848225923244026
```

```
!pip install my-package
```

```
Collecting my-package
  Downloading my_package-0.0.0-py3-none-any.whl (2.0 kB)
Installing collected packages: my-package
Successfully installed my-package-0.0.0
```

```
!pip install pydotplus
```

```
Requirement already satisfied: pydotplus in /usr/local/lib/python3.10/dist-packages (2.0.2)
Requirement already satisfied: pyparsing>=2.0.1 in /usr/local/lib/python3.10/dist-packages (from pydotplus) (3.1.1)
```

```
# Importing required packages for visualization
from IPython.display import Image
from six import StringIO
from sklearn.tree import export_graphviz
import pydotplus,graphviz
```

```
# Putting features
```

```
features = list(df.columns[1:])
```

```
features
```

```
['fnlwgt',
 'educational-num',
 'capital-gain',
 'capital-loss',
 'hours-per-week',
 'workclass',
 'education',
 'marital-status',
 'occupation',
 'relationship',
 'race',
 'gender',
 'native-country',
 'income']
```

```
!pip install graphviz
```

Requirement already satisfied: graphviz in /usr/local/lib/python3.10/dist-packages (0.20.1)

```
# plotting tree with max_depth=3
```


```
dot_data = StringIO()
```

```
export_graphviz(dt_default, out_file=dot_data,
```

```
feature_names=features, filled=True,rounded=True)
```

```
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
```

```
Image(graph.create_png())
```



```
# GridSearchCV to find optimal max_depth
```

```
from sklearn.model_selection import KFold
```

```
from sklearn.model_selection import GridSearchCV
```

```
# specify number of folds for k-fold CV
```

```
n_folds = 5
```

```
# parameters to build the model on
```

```
parameters = {'max_depth': range(1, 40)}
```

```
# instantiate the model
```

```
dtree = DecisionTreeClassifier(criterion = "gini",
```

```
random_state = 100)
```


```
# fit tree on training data
```

```
tree = GridSearchCV(dtree, parameters,
```

```
cv=n_folds,
```

```
scoring="accuracy")
```


```
tree.fit(X_train, y_train)
```



```
# scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()
```

| | mean_fit_time | std_fit_time | mean_score_time | std_score_time | param_max_depth | |
|---|---------------|--------------|-----------------|----------------|-----------------|--------|
| 0 | 0.017508 | 0.001094 | 0.003653 | 0.000255 | 1 | {'max_ |
| 1 | 0.026208 | 0.000638 | 0.003497 | 0.000050 | 2 | {'max_ |
| 2 | 0.035516 | 0.000600 | 0.003571 | 0.000109 | 3 | {'max_ |
| 3 | 0.046339 | 0.002691 | 0.003698 | 0.000265 | 4 | {'max_ |
| 4 | 0.054046 | 0.000881 | 0.003738 | 0.000126 | 5 | {'max_ |


```
# GridSearchCV to find optimal max_depth
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
# specify number of folds for k-fold CV
n_folds = 5
# parameters to build the model on
parameters = {'min_samples_leaf': range(5, 200, 20)}
# instantiate the model
dtree = DecisionTreeClassifier(criterion = "gini",
random_state = 100)
# fit tree on training data
tree = GridSearchCV(dtree, parameters,
cv=n_folds,
scoring="accuracy")
tree.fit(X_train, y_train)
```



```
# scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()
```

| | mean_fit_time | std_fit_time | mean_score_time | std_score_time | param_min_samples_leaf | |
|---|---------------|--------------|-----------------|----------------|------------------------|--|
| 0 | 0.137336 | 0.004291 | 0.004951 | 0.000940 | 5 | |
| 1 | 0.115085 | 0.003973 | 0.004306 | 0.000304 | 25 | |
| 2 | 0.107868 | 0.005090 | 0.004147 | 0.000113 | 45 | |
| 3 | 0.101131 | 0.002170 | 0.004166 | 0.000221 | 65 | |
| 4 | 0.100600 | 0.004742 | 0.004072 | 0.000051 | 85 | |

```
# GridSearchCV to find optimal min_samples_split
from sklearn.model_selection import KFold
from sklearn.model_selection import GridSearchCV
# specify number of folds for k-fold CV
n_folds = 5
# parameters to build the model on
parameters = {'min_samples_split': range(5, 200, 20)}
# instantiate the model
dtree = DecisionTreeClassifier(criterion = "gini",
random_state = 100)
# fit tree on training data
tree = GridSearchCV(dtree, parameters,
cv=n_folds,
scoring="accuracy")
tree.fit(X_train, y_train)
```



```
# scores of GridSearch CV
scores = tree.cv_results_
pd.DataFrame(scores).head()
```

| | mean_fit_time | std_fit_time | mean_score_time | std_score_time | param_min_samples_spli |
|---|---------------|--------------|-----------------|----------------|------------------------|
| 0 | 0.153947 | 0.002147 | 0.005929 | 0.000578 | |
| 1 | 0.143463 | 0.004214 | 0.005008 | 0.000071 | 2 |
| 2 | 0.138373 | 0.001752 | 0.005027 | 0.000082 | 4 |
| 3 | 0.133974 | 0.005993 | 0.004872 | 0.000038 | 6 |
| 4 | 0.133274 | 0.005639 | 0.005146 | 0.000392 | 8 |

```
# Create the parameter grid
param_grid = {
'max_depth': range(5, 15, 5),
'min_samples_leaf': range(50, 150, 50),
'min_samples_split': range(50, 150, 50),
'criterion': ["entropy", "gini"]
}
n_folds = 5

# Instantiate the grid search model
dtree = DecisionTreeClassifier()
grid_search = GridSearchCV(estimator = dtree, param_grid = param_grid,
cv = n_folds, verbose = 1)
# Fit the grid search to the data
grid_search.fit(X_train,y_train)
```

Fitting 5 folds for each of 16 candidates, totalling 80 fits

```
GridSearchCV
```

```
# cv results
cv_results = pd.DataFrame(grid_search.cv_results_)
cv_results
```

| | | | | | |
|---|----------|----------|----------|----------|---------|
| 6 | 0.137194 | 0.005227 | 0.006264 | 0.001309 | entropy |
| 7 | 0.122113 | 0.018699 | 0.004998 | 0.000869 | entropy |
| 8 | 0.052541 | 0.002334 | 0.003601 | 0.000090 | gini |
| 9 | 0.051650 | 0.000660 | 0.003619 | 0.000087 | gini |

```
# printing the optimal accuracy score and hyperparameters
print("best accuracy", grid_search.best_score_)
print(grid_search.best_estimator_)
```

```
best accuracy 0.8523105983446813
DecisionTreeClassifier(max_depth=10, min_samples_leaf=50, min_samples_split=50)
```


```
# model with optimal hyperparameters
clf_gini = DecisionTreeClassifier(criterion = "gini",
random_state = 100,
max_depth=10,
min_samples_leaf=50,
min_samples_split=50)
clf_gini.fit(X_train, y_train)
```

```
DecisionTreeClassifier
DecisionTreeClassifier(max_depth=10, min_samples_leaf=50, min_samples_split=50,
random_state=100)
```

```
# accuracy score
clf_gini.score(X_test,y_test)
```

```
0.852860246198407
```

```
#plotting the tree
dot_data = StringIO()
export_graphviz(clf_gini, out_file=dot_data,feature_names=features,filled=True,rounded=True)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
Image(graph.create_png())
```



```
# tree with max_depth = 3
clf_gini = DecisionTreeClassifier(criterion = "gini",
random_state = 100,
max_depth=3,
min_samples_leaf=50,
min samples split=50)
```



```

clf_gini.fit(X_train, y_train)
# score
print(clf_gini.score(X_test,y_test))
# plotting tree with max_depth=3
dot_data = StringIO()
export_graphviz(clf_gini, out_file=dot_data,feature_names=features,filled=True,rounded=True)

```

```
0.8331643736422882
```

```

# classification metrics
from sklearn.metrics import classification_report,confusion_matrix
y_pred = clf_gini.predict(X_test)
print(classification_report(y_test, y_pred))
# confusion matrix
print(confusion_matrix(y_test,y_pred))

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.84 | 0.95 | 0.90 | 10341 |
| 1 | 0.77 | 0.47 | 0.59 | 3469 |
| accuracy | | | 0.83 | 13810 |
| macro avg | 0.81 | 0.71 | 0.74 | 13810 |
| weighted avg | 0.83 | 0.83 | 0.82 | 13810 |

```

[[9860 481]
 [1823 1646]]

```