

```
# Import libraries
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
sns.set(style='white', context='notebook', palette='deep')
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV, cross_val_score, StratifiedKFold, learning_curve, train_test_split, KF
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score

# To ignore warning messages
import warnings
warnings.filterwarnings('ignore')
# Adult dataset path
adult_dataset_path = "/content/adult.csv"
# Function for loading adult dataset
def load_adult_data(adult_path=adult_dataset_path):
    csv_path = os.path.join(adult_path)
    return pd.read_csv(csv_path)

# Calling load adult function and assigning to a new variable df
df = load_adult_data()
# load top 3 rows values from adult dataset
df.head(3)
```

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relation
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Hus
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Hus

```
print ("Rows : " ,df.shape[0])
print ("Columns : " ,df.shape[1])
print ("\nFeatures : \n" ,df.columns.tolist())
print ("\nMissing values : " , df.isnull().sum().values.sum())
print ("\nUnique values : \n",df.nunique())
```

```
Rows : 48842
Columns : 15
```

```
Features :
['age', 'workclass', 'fnlwgt', 'education', 'educational-num', 'marital-status', 'occupation', 'relationship', 'r
```

```
Missing values : 0
```

```
Unique values :
age          74
workclass     9
fnlwgt      28523
```

```
education          16
educational-num    16
marital-status      7
occupation         15
relationship        6
race               5
gender             2
capital-gain       123
capital-loss       99
hours-per-week     96
native-country     42
income            2
dtype: int64
```

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48842 entries, 0 to 48841
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   48842 non-null  int64
1   workclass             48842 non-null  object
2   fnlwgt               48842 non-null  int64
3   education             48842 non-null  object
4   educational-num       48842 non-null  int64
5   marital-status       48842 non-null  object
6   occupation            48842 non-null  object
7   relationship          48842 non-null  object
8   race                 48842 non-null  object
9   gender               48842 non-null  object
10  capital-gain          48842 non-null  int64
11  capital-loss          48842 non-null  int64
12  hours-per-week        48842 non-null  int64
13  native-country        48842 non-null  object
14  income                48842 non-null  object
dtypes: int64(6), object(9)
memory usage: 5.6+ MB
```

df.describe()

	age	fnlwgt	educational-num	capital-gain	capital-loss	hours
count	48842.000000	4.884200e+04	48842.000000	48842.000000	48842.000000	48842.0
mean	38.643585	1.896641e+05	10.078089	1079.067626	87.502314	40.4
std	13.710510	1.056040e+05	2.570973	7452.019058	403.004552	12.3
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1.0
25%	28.000000	1.175505e+05	9.000000	0.000000	0.000000	40.0
50%	37.000000	1.781445e+05	10.000000	0.000000	0.000000	40.0
75%	48.000000	2.376420e+05	12.000000	0.000000	0.000000	45.0

df.head()

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relation
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Hus
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Hus

```
# checking "?" total values present in particular 'workclass' feature
```

```
df_check_missing_workclass = (df['workclass']=='?').sum()
```

```
df_check_missing_workclass
```

```
2799
```

```
# checking "?" total values present in particular 'occupation' feature
```

```
df_check_missing_occupation = (df['occupation']=='?').sum()
```

```
df_check_missing_occupation
```

```
2809
```

```
# checking "?" values, how many are there in the whole dataset
```

```
df_missing = (df=='?').sum()
```

```
df_missing
```

```
age                0
workclass          2799
fnlwgt             0
education          0
educational-num    0
marital-status     0
occupation         2809
relationship       0
race               0
gender             0
capital-gain       0
capital-loss       0
hours-per-week     0
native-country     857
income             0
dtype: int64
```

```
percent_missing = (df=='?').sum() * 100/len(df)
```

```
percent_missing
```

```
age                0.000000
workclass          5.730724
fnlwgt             0.000000
education          0.000000
educational-num    0.000000
marital-status     0.000000
occupation         5.751198
relationship       0.000000
race               0.000000
gender             0.000000
capital-gain       0.000000
capital-loss       0.000000
hours-per-week     0.000000
native-country     1.754637
income             0.000000
dtype: float64
```

```
# find total number of rows which doesn't contain any missing value as '?'
df.apply(lambda x: x != '?',axis=1).sum()
```

```
age          48842
workclass    46043
fnlwt        48842
education    48842
educational-num 48842
marital-status 48842
occupation   46033
relationship 48842
race         48842
gender       48842
capital-gain 48842
capital-loss 48842
hours-per-week 48842
native-country 47985
income       48842
dtype: int64
```

```
# dropping the rows having missing values in workclass
df = df[df['workclass'] != '?']
df.head()
```

	age	workclass	fnlwt	education	educational-num	marital-status	occupation	relation
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Hus
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Hus
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Hus
5	34	Private	198693	10th	6	Never-married	Other-service	Not-in-fa

```
# select all categorical variables
df_categorical = df.select_dtypes(include=['object'])
# checking whether any other column contains '?' value
df_categorical.apply(lambda x: x == '?',axis=1).sum()
```

```
workclass      0
education      0
marital-status 0
occupation     10
relationship    0
race           0
gender         0
native-country 811
income         0
dtype: int64
```

```
from sklearn import preprocessing
# encode categorical variables using label Encoder
# select all categorical variables
df_categorical = df.select_dtypes(include=['object'])
```

```
df_categorical.head()
```

	workclass	education	marital-status	occupation	relationship	race	gender	native-country
0	Private	11th	Never-married	Machine-op-inspct	Own-child	Black	Male	United-States
1	Private	HS-grad	Married-civ-spouse	Farming-fishing	Husband	White	Male	United-States
			Married-					

```
# apply label encoder to df_categorical
le = preprocessing.LabelEncoder()
df_categorical = df_categorical.apply(le.fit_transform)
df_categorical.head()
```

	workclass	education	marital-status	occupation	relationship	race	gender	native-country
0	2	1	4	6	3	2	1	39
1	2	11	2	4	0	4	1	39
2	1	7	2	10	0	4	1	39
3	2	15	2	6	0	2	1	39

```
# Next, Concatenate df_categorical dataframe with original df (dataframe)
# first, Drop earlier duplicate columns which had categorical values
df = df.drop(df_categorical.columns,axis=1)
df = pd.concat([df,df_categorical],axis=1)
df.head()
```

	age	fnlwgt	educational-num	capital-gain	capital-loss	hours-per-week	workclass	education	marital-status
0	25	226802	7	0	0	40	2	1	4
1	38	89814	9	0	0	50	2	11	2
2	28	336951	12	0	0	40	1	7	2
3	44	160323	10	7688	0	40	2	15	2
5	34	198693	6	0	0	30	2	0	4

```
# look at column type
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 46033 entries, 0 to 48841
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   46033 non-null  int64
1   fnlwgt                46033 non-null  int64
2   educational-num       46033 non-null  int64
3   capital-gain          46033 non-null  int64
4   capital-loss          46033 non-null  int64
5   hours-per-week        46033 non-null  int64
6   workclass             46033 non-null  int64
```

```

7 education      46033 non-null int64
8 marital-status  46033 non-null int64
9 occupation     46033 non-null int64
10 relationship   46033 non-null int64
11 race          46033 non-null int64
12 gender        46033 non-null int64
13 native-country 46033 non-null int64
14 income        46033 non-null int64

```


```
dtypes: int64(15)
```

```
memory usage: 5.6 MB
```

```

plt.figure(figsize=(14,10))
sns.heatmap(df.corr(),annot=True,fmt='.2f')
plt.show()

```



```

# convert target variable income to categorical
df['income'] = df['income'].astype('category')
# check df info again whether everything is in right format or not
df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 46033 entries, 0 to 48841
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -

```

```

0  age                46033 non-null int64
1  fnlwgt             46033 non-null int64
2  educational-num    46033 non-null int64
3  capital-gain       46033 non-null int64
4  capital-loss       46033 non-null int64
5  hours-per-week     46033 non-null int64
6  workclass          46033 non-null int64
7  education          46033 non-null int64
8  marital-status     46033 non-null int64
9  occupation         46033 non-null int64
10 relationship       46033 non-null int64
11 race              46033 non-null int64
12 gender            46033 non-null int64
13 native-country    46033 non-null int64
14 income            46033 non-null category
dtypes: category(1), int64(14)
memory usage: 5.3 MB

```

```

# Importing train_test_split
from sklearn.model_selection import train_test_split
# Putting independent variables/features to X
X = df.drop('income',axis=1)
# Putting response/dependent variable/feature to y
y = df['income']
X.head(3)

```

	age	fnlwgt	educational-num	capital-gain	capital-loss	hours-per-week	workclass	education	marital-status
0	25	226802	7	0	0	40	2	1	
1	38	89814	9	0	0	50	2	11	
2	28	336951	12	0	0	40	1	7	

```
y.head(3)
```

```

0    0
1    0
2    1
Name: income, dtype: category
Categories (2, int64): [0, 1]

```

```

# Splitting the data into train and test
X_train,X_test,y_train,y_test = train_test_split(X,y)
X_train.head()

```

	age	fnlwgt	educational-num	capital-gain	capital-loss	hours-per-week	workclass	education	marital-status
13554	58	196502	10	0	0	60	2	15	
46282	27	297457	9	0	0	40	2	11	
25679	27	30244	9	0	0	80	4	11	
8775	42	165309	9	0	0	50	2	11	
13211	68	140892	14	0	0	15	4	12	

```

test_size = 0.20
seed = 7
num_folds = 10

```

```

scoring = 'accuracy'
# Params for Random Forest
num_trees = 100
max_features = 3

random_forest = RandomForestClassifier(n_estimators=250,max_features=5)
random_forest.fit(X_train, y_train)
predictions = random_forest.predict(X_test)
print("Accuracy: %s%%" % (100*accuracy_score(y_test, predictions)))
print(confusion_matrix(y_test, predictions))
print(classification_report(y_test, predictions))

```

Accuracy: 85.44617256060475%

[[8015 628]

[1047 1819]]

	precision	recall	f1-score	support
0	0.88	0.93	0.91	8643
1	0.74	0.63	0.68	2866
accuracy			0.85	11509
macro avg	0.81	0.78	0.80	11509
weighted avg	0.85	0.85	0.85	11509



Experiment No. 4
Apply Random Forest Algorithm on Adult Census Income Dataset and analyze the performance of the model
Date of Performance:14-08-23
Date of Submission:22-08-23



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

Aim: Apply Random Forest Algorithm on Adult Census Income Dataset and analyze the performance of the model.

Objective: Able to perform various feature engineering tasks, apply Random Forest Algorithm on the given dataset and maximize the accuracy, Precision, Recall, F1 score.

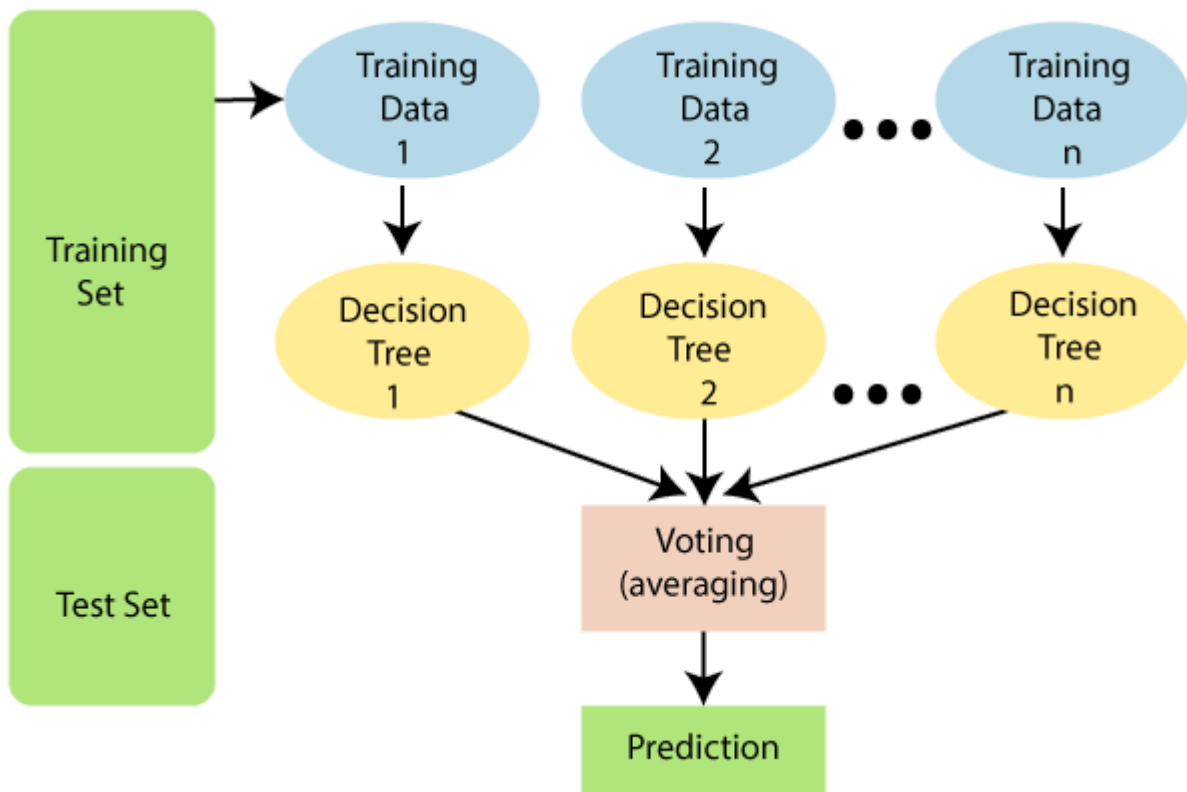
Theory:

Random Forest is a popular machine learning algorithm that belongs to the supervised learning technique. It can be used for both Classification and Regression problems in ML. It is based on the concept of ensemble learning, which is a process of combining multiple classifiers to solve a complex problem and to improve the performance of the model.

As the name suggests, "Random Forest is a classifier that contains a number of decision trees on various subsets of the given dataset and takes the average to improve the predictive accuracy of that dataset." Instead of relying on one decision tree, the random forest takes the prediction from each tree and based on the majority votes of predictions, and it predicts the final output.

The greater number of trees in the forest leads to higher accuracy and prevents the problem of overfitting.

The below diagram explains the working of the Random Forest algorithm:



Dataset:

Predict whether income exceeds \$50K/yr based on census data. Also known as "Adult" dataset.

Attribute Information:

Listing of attributes:

>50K, <=50K.

age: continuous.

workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.

fnlwgt: continuous.



Vidyavardhini's College of Engineering & Technology

Department of Computer Engineering

education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.

education-num: continuous.

marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.

occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.

relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.

race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.

sex: Female, Male.

capital-gain: continuous.

capital-loss: continuous.

hours-per-week: continuous.

native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad & Tobago, Peru, Hong, Holand-Netherlands.

Code:



Conclusion:

The correlations among these variables are generally simpler in strength, they lack strong linear associations with one another.

Age exhibits a weak positive correlation with both education number and hours worked per week. Education numbers display a light positive correlation with capital gains. Also, there exists a weak negative correlation between capital gains and capital losses.

Accuracy: The model's accuracy is 85.44%, correctly predicting income levels for most instances.

Confusion Matrix: True positives (8015), False positives (628), and False negatives (1047), True negatives (1819) predictions.

Precision: Precision for income 0 = 0.08 and precision for income 1 = 0.74

Recall: Recall for income 0 = 0.93 and recall for income 1 = 0.63

F1-score: F1-score is the mean between precision and recall, indicating overall model effectiveness. It contains 0 for 0.91 and 1 for 0.68

Random Forest tends to provide better results than a Decision Tree. The Random Forest model combines the predictions of multiple Decision Trees, which can lead to improved accuracy and generalization.